First Steps to C Mastery

Today's Lecture

First steps towards C mastery

- ■ Arithmetic: adders, wraparound

- ■ Software: hallmarks of good software

- ■ Pointers, structs, and memory

- ■ More on Heap Allocator

But What Type Is It?

```
unsigned int b = 2;        int b = 2;
int a = -2;                int a = -2;

if(a > b)                  if(a > b)
   printf("a > b\n");         printf("a > b\n");
else                       else
   printf("b > a\n");         printf("b > a\n");
```

But What Type Is It?

```
unsigned int b = 2;        int b = 2;
int a = -2;                int a = -2;

if(a > b)                  if(a > b)
   printf("a > b\n");         printf("a > b\n");
else                       else
   printf("b > a\n");         printf("b > a\n");
```

C converts a into unsigned

Conversions Everywhere

```c
#include <stdio.h>
#include <limits.h>
#include <assert.h>

int main(void) {
  assert(sizeof(unsigned char)==1);

  unsigned char uc1 = 0xff;
  unsigned char uc2 = 0;

  if(~uc1 == uc2) {
    printf("%hhx == %hhx\n", ~uc1, uc2);
  } else {
    printf("%hhx != %hhx\n", ~uc1, uc2);
  }
  return 0;
}
```

Operand Conversion

	u8	u16	u32	u64	i8	i16	i32	i64
u8	i32	i32	u32	u64	i32	i32	i32	i64
u16	i32	i32	u32	u64	i32	i32	i32	i64
u32	u32	u32	u32	u64	u32	u32	u32	i64
u64	u64	u64	u64	u64	u64	u64	u64	u64
i8	i32	i32	u32	u64	i32	i32	i32	i64
i16	i32	i32	u32	u64	i32	i32	i32	i64
i32	i32	i32	u32	u64	i32	i32	i32	i64
i64	i64	i64	i64	u64	i64	i64	i64	i64

code/wraparound

- In two's complement, when you exceed the maximum value of int (2,147,483,647), you "wrap around" to negative numbers:

PSY - GANGNAM STYLE (강남스타일) M/V
officialpsy
Subscribe 7,600,830
-2142584554
8,761,309   1,139,933
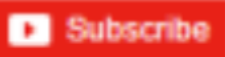Add to   Share   ••• More

- Here is the link after Google upgraded to 64-bit integers:

PSY - GANGNAM STYLE (강남스타일) M/V
officialpsy
Subscribe 7,600,830
2,152,382,740
8,761,309   1,139,933
Add to   Share   ••• More

Writing Good Systems Software

```c
void serial_init() {
    unsigned int ra;

    // Configure the UART
    PUT32(AUX_ENABLES, 1);
    PUT32(AUX_MU_IER_REG, 0); // Clear FIFO
    PUT32(AUX_MU_CNTL_REG, 0); // Default RTS/CTS
    PUT32(AUX_MU_LCR_REG, 3); // Put in 8 bit mode
    PUT32(AUX_MU_MCR_REG, 0); // Default RTS/CTS auto flow control
    PUT32(AUX_MU_IER_REG, 0); // Clear FIFO
    PUT32(AUX_MU_IIR_REG, 0xC6); // Baudrate
    PUT32(AUX_MU_BAUD_REG, 270); // Baudrate

    // Configure the GPIO lines
    ra = GET32(GPFSEL1);
    ra &= ~(7 << 12); //gpio14
    ra |= 2 << 12;     //alt5
    ra &= ~(7 << 15); //gpio15
    ra |= 2 << 15;     //alt5
    PUT32(GPFSEL1,ra);
    PUT32(GPPUD,0);
    for (ra = 0; ra < 150; ra++) dummy(ra);
    PUT32(GPPUDCLK0, (1 << 14) | (1 << 15));
    for (ra = 0; ra < 150; ra++) dummy(ra);
    PUT32(GPPUDCLK0, 0);

    // Enable the serial port (both RX and TX)
    PUT32(AUX_MU_CNTL_REG, 3);
}
```

```c
void uart_init(void) {
    int *aux = (int*)AUX_ENABLES;;

    *aux = AUX_ENABLE; // turn on mini-uart

    uart->ier  = 0;
    uart->cntl = 0;
    uart->lcr  = MINI_UART_LCR_8BIT;
    uart->mcr  = 0;
    uart->ier  = 0;
    uart->iir  = MINI_UART_IIR_RX_FIFO_CLEAR |
                 MINI_UART_IIR_RX_FIFO_ENABLE |
                 MINI_UART_IIR_TX_FIFO_CLEAR |
                 MINI_UART_IIR_TX_FIFO_ENABLE;

    uart->baud =  270; // baud rate ((250,000,000/115200)/8)-1 = 270

    gpio_set_function(GPIO_TX, GPIO_FUNC_ALT5);
    gpio_set_function(GPIO_RX, GPIO_FUNC_ALT5);

    uart->cntl = MINI_UART_CNTL_TX_ENABLE |
                 MINI_UART_CNTL_RX_ENABLE;
}
```

Well-written software is easy for someone to read and understand.

Well-written software is easy for someone to read and understand.

Code that is easier to understand has fewer bugs.

Long comments != easy to read and understand.

Understand at the line, function, file, and structural levels.

Lesson: Imagine someone else has to fix a bug in your code: what should it look like make this easier? *Hint: be a section leader, you'll have to read student code and you'll learn a lot.*

Systems Code Is Terse But Unforgiving

Systems Code Is Terse But Unforgiving

Next week in lab you will write a "PS/2 scan code reader" to read keyboard input: if any part of it is wrong, you won't read scan codes. It's only 20-30 lines of code!

The mailbox code you'll use for the frame buffer is ~10 lines of code: we once debugged it for 9 hours.

Lesson: if you know exactly what you have to do it can take only minutes; throwing away and rewriting can often be *faster*. Sunk cost fallacy!

```c
void uart_init(void) {
    int *aux = (int*)AUX_ENABLES;;

    *aux = AUX_ENABLE; // turn on mini-uart

    uart->ier  = 0;
    uart->cntl = 0;
    uart->lcr  = MINI_UART_LCR_8BIT;
    uart->mcr  = 0;
    uart->ier  = 0;
    uart->iir  = MINI_UART_IIR_RX_FIFO_CLEAR |
                 MINI_UART_IIR_RX_FIFO_ENABLE |
                 MINI_UART_IIR_TX_FIFO_CLEAR |
                 MINI_UART_IIR_TX_FIFO_ENABLE;

    uart->baud =  270; // baud rate ((250,000,000/115200)/8)-1 = 270

    gpio_set_function(GPIO_TX, GPIO_FUNC_ALT5);
    gpio_set_function(GPIO_RX, GPIO_FUNC_ALT5);

    uart->cntl = MINI_UART_CNTL_TX_ENABLE | MINI_UART_CNTL_RX_ENABLE;;
}
```

Pointers, Structures, Etc.

Computer Memory

0x0200000000
(8GB)

0x0000000000

one byte (8 bits) each

0x0   0x20

0x5e0   0x5ff
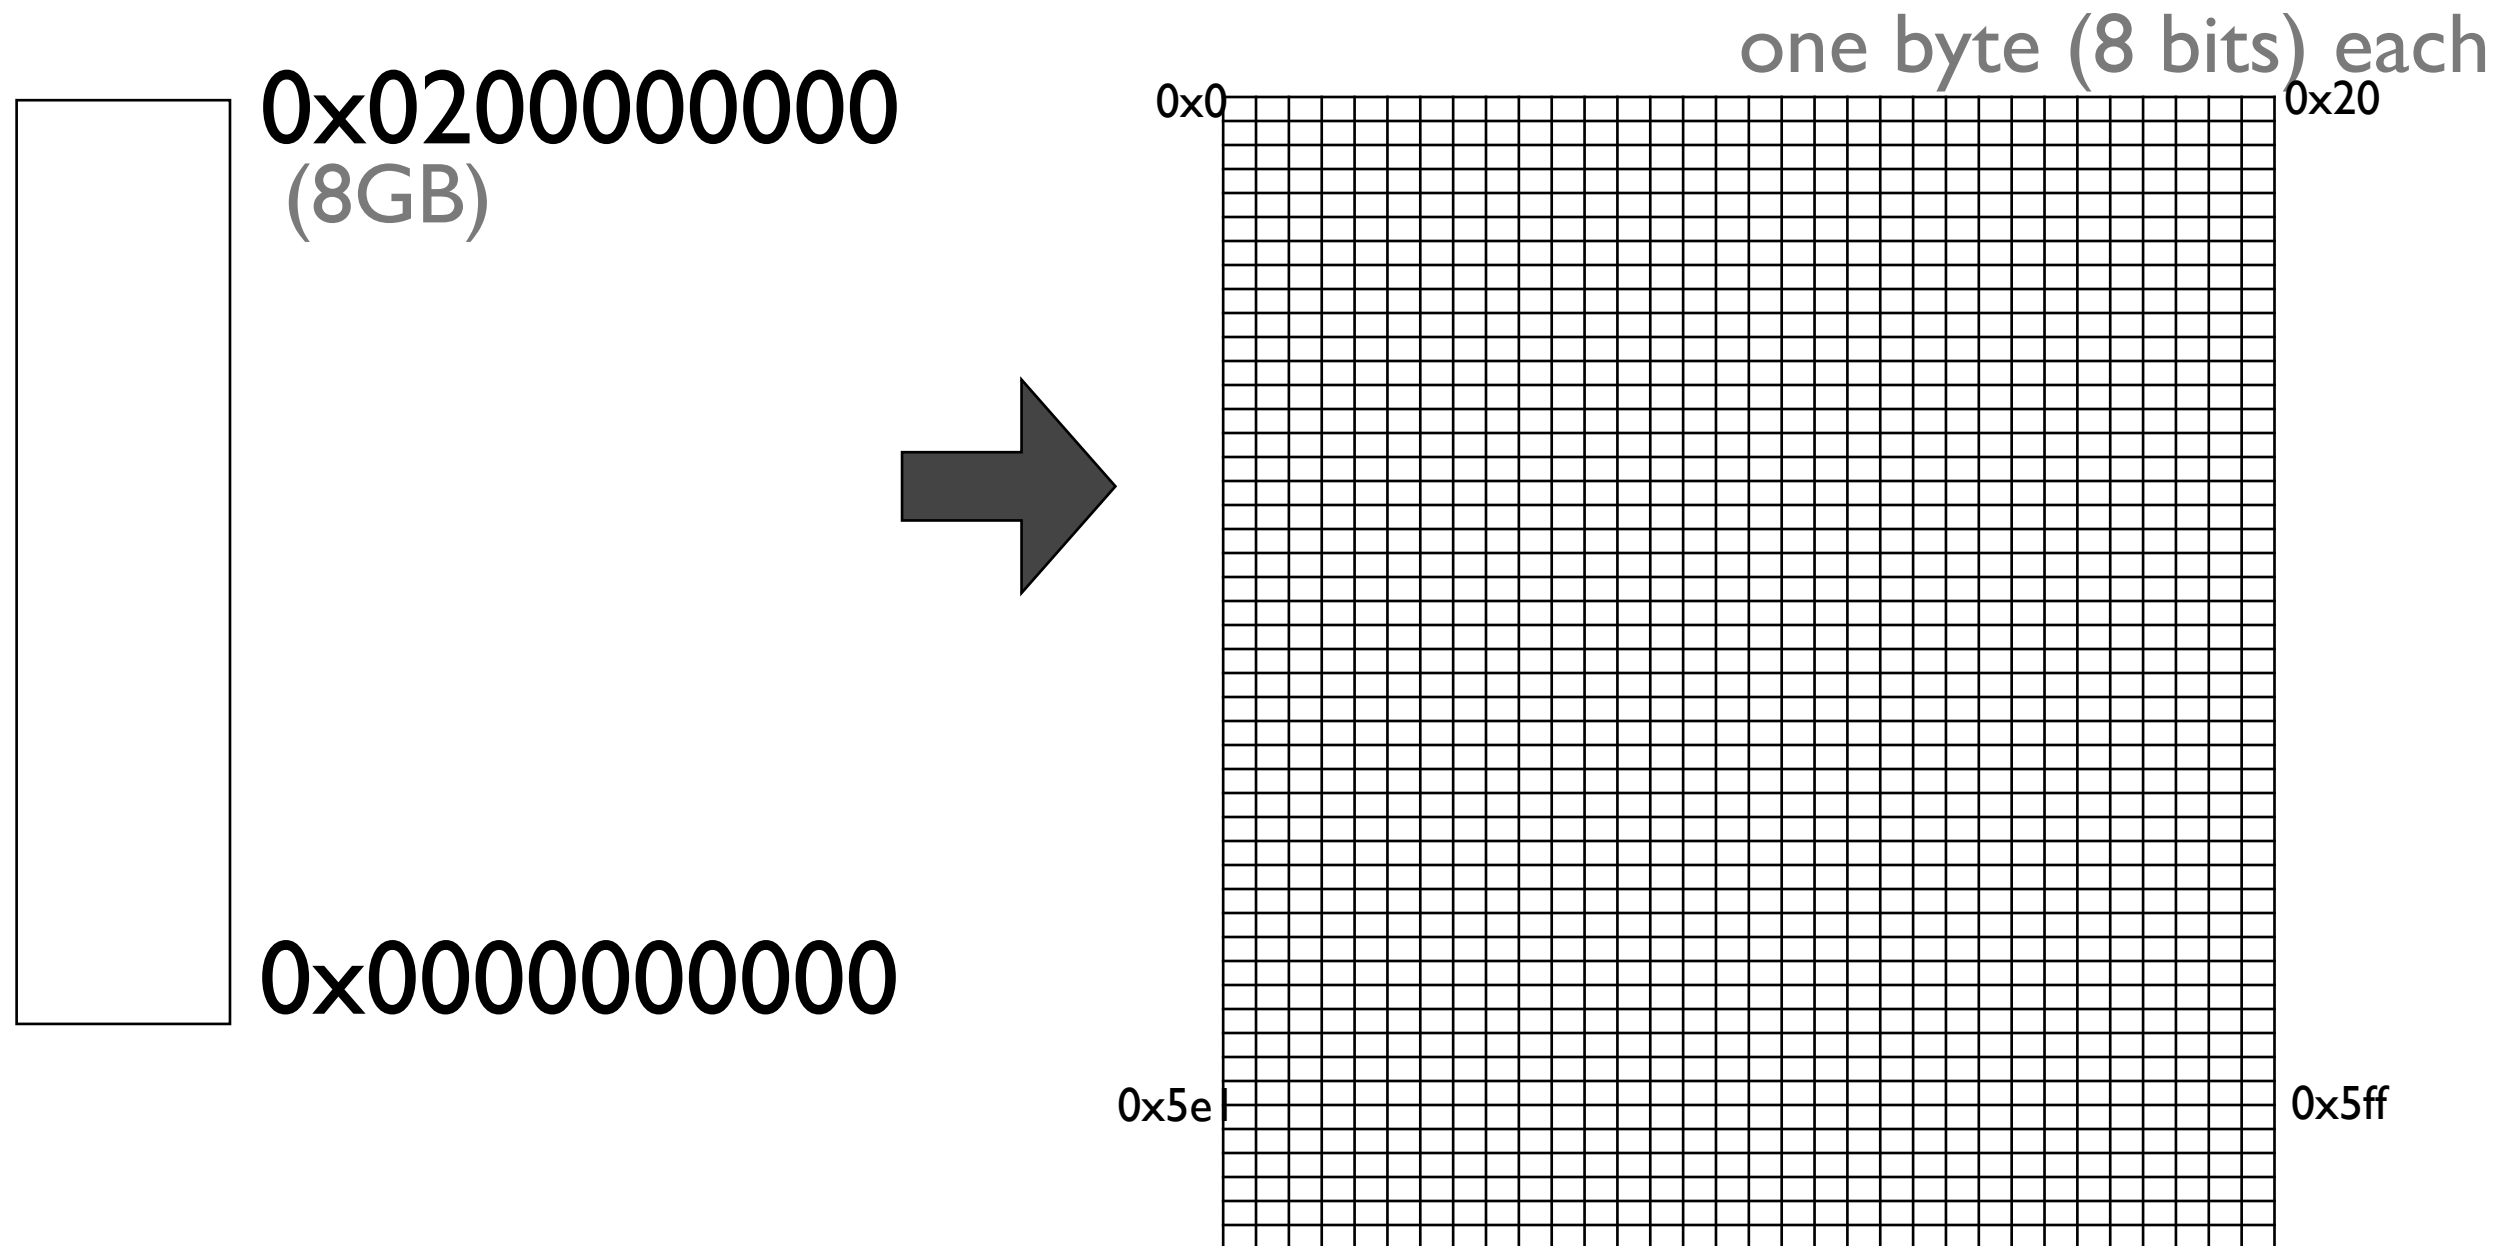
Endianness

Multibyte words: how do you arrange the bytes?

$$1{,}024 = 0x0400 =$$

?	?

Little endian: least significant byte is at lowest address

■ Makes most sense from an addressing/computational standpoint

0x00	0x04

Big endian: most significant byte is at lowest address

■ Makes most sense to a human reader

0x04	0x00

"Arrays are Pointers"

Sort of true. Sometimes.

When are arrays and pointers different?

- A pointer is a location in memory storing an address (you can change the pointer)

- An array is a location in memory storing data (you can change the *elements* of an array, but not the array itself)

code/pointers

C structs

```
struct data {
  unsigned char fields;
  unsigned int num_changes;
  char name[MAX_NAME + 1];
  unsigned short references;
  unsigned short links;
}
```

How big is this structure?
How is it laid out in memory (on an ARM)?

C structs

A C struct is just a shorthand and convenient way to allocate memory and describe offsets from a pointer.

```c
struct data {
  unsigned char fields;
  unsigned int num_changes;
  char name[4];
  unsigned short references;
  unsigned short links;
};

struct data* d = …;
d->fields = 0;          // d + 0
d->num_changes = 0;  // d + 4
d->references = 0;    // d + 12
d->links = 0;          // d + 14
```

Structures

Example:

```
struct rec {
    int i;
    int j;
    int a[2];
    int *p;
};
```

This structure has four fields: two 4-byte values of type `int`, a two-element array of type `int`, and an 4-byte `int` pointer, for a total of 20 bytes:

Offset	0	4	8		16	20

Contents	i	j	a[0]	a[1]	p

- The numbers along the top of the diagram are the byte offsets of the fields from the beginning of the structure.
- Struct layout is guaranteed to be in the order you defined, however the exact location is dependent on alignment in the memory system (ARM requires 4-byte alignment).
- Note that the array is embedded in the structure.
- To access the fields, the compiler generates code that adds the field offset to the address of the structure.

Structures

Example:

```
struct rec {
    int i;
    int j;
    int a[2];
    int *p;
};
```

This structure has four fields: two 4-byte values of type `int`, a two-element array of type `int`, and an 4-byte `int` pointer, for a total of 20 bytes:

Offset	0	4	8		16	20

Contents	i	j	a[0]	a[1]	p

- Let's look at the following code in gdb:

```
struct rec r;
r.i = 1;
r.j = 2;
r.a[0] = 3;
r.a[1] = 4;
r.p = (int *)0x8000;
```

What's this about alignment?

Example:

Let's look at a different struct.

```
struct withchar {
    int i;
    int j;
    char a;
    char b;
    int k;
};
```

Offset	0	4	8	12	16

Contents	i	j	ab	p

- It looks like there is some waste! There are two unused bytes at locations 10 and 11 -- this is because in order to meet alignment requirements, the compiler had to put k at location 12, and it couldn't put it at location 10 (not divisible by 4).

- Let's look at the following code in gdb:

```
    struct withchar wc;
wc.i = 1;
wc.j = 2;
wc.a = 'a';
wc.b = 'b';
wc.k = 3;
```

- Remember: the compiler *must* put your struct in order, but it also must meet alignment requirements.

Laying a struct onto memory

```
struct rec {
    int i;
    int j;
    int a[2];
    int *p;
};

void main(void)
{
    char arr[1024];
    for (int i=0; i < 1024 / 4; i++) {
        *((int *)arr + i) = 0xdeadbeef; // fill with deadbeef
    }
    char *randomLocation = &arr[64];
    ((struct rec *)randomLocation)->i = 5;
    ((struct rec *)randomLocation)->j = 6;
    ((struct rec *)randomLocation)->a[0] = 7;
    ((struct rec *)randomLocation)->a[1] = 8;
    ((struct rec *)randomLocation)->p = (int *)0x8000;
```

- We can put a struct anywhere, as long as we align it to a 4-byte boundary

Laying a struct onto memory

```
(gdb) x/40x arr
0x7fffbb4:      0xdeadbeef      0xdeadbeef      0xdeadbeef      0xdeadbeef
0x7fffbc4:      0xdeadbeef      0xdeadbeef      0xdeadbeef      0xdeadbeef
0x7fffbd4:      0xdeadbeef      0xdeadbeef      0xdeadbeef      0xdeadbeef
0x7fffbe4:      0xdeadbeef      0xdeadbeef      0xdeadbeef      0xdeadbeef
0x7fffbf4:      0x00000005      0x00000006      0x00000007      0x00000008
0x7fffc04:      0x00008000      0xdeadbeef      0xdeadbeef      0xdeadbeef
0x7fffc14:      0xdeadbeef      0xdeadbeef      0xdeadbeef      0xdeadbeef
0x7fffc24:      0xdeadbeef      0xdeadbeef      0xdeadbeef      0xdeadbeef
0x7fffc34:      0xdeadbeef      0xdeadbeef      0xdeadbeef      0xdeadbeef
0x7fffc44:      0xdeadbeef      0xdeadbeef      0xdeadbeef      0xdeadbeef
(gdb)
```

- On Monday, we discussed the *bump allocator*, which creates a dumb malloc. For assign4, you will implement an *implicit free list*, which is a better way to create a heap allocator. It uses what is called a **block header** to hold the information. In the assignment diagrams page (http://cs107e.github.io/assignments/assign4/images/block_hdr/) there is a potential header defined as follows*:

```
struct header {
    size_t payload_size;
    int status;           // 0 if free, 1 if in use
};
```

- The block header is actually *stored in the same memory area as the payload, and it generally precedes the payload.*
- Let's take a look at a series of mallocs, frees, and reallocs and see what happens in this heap.
- We must align each pointer returned to the malloc caller on a 4-byte boundary (even if we will waste some bytes).

*feel free to define a header more efficiently!

96 bytes

00	04	08	0c	10	14	18	1c	20	24	28	2c	30	34	38	3c	40	44	48	4c	50	54	58	5c
?	?	?	?	?	?	?	?	?	?	?	?	?	?	?	?	?	?	?	?	?	?	?	?
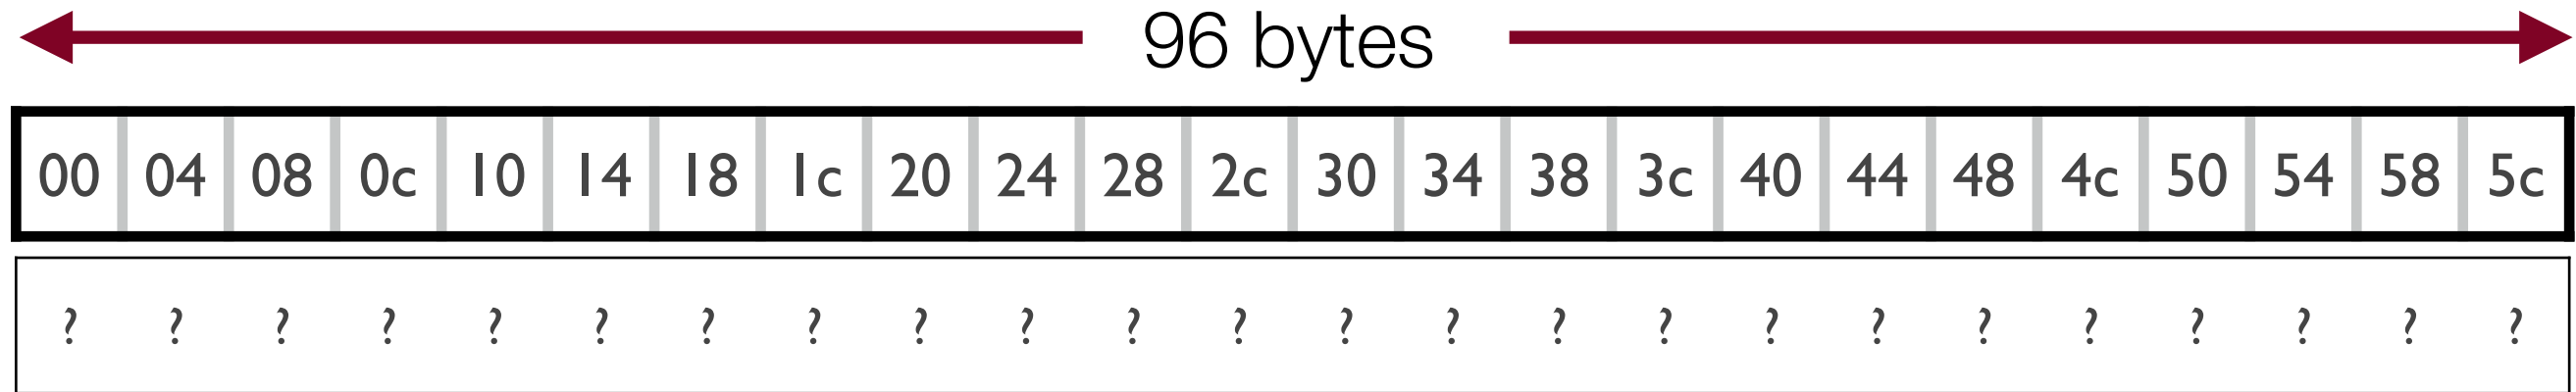
- When the heap is initialized, assume there are 96 bytes starting at location 0x9000. Note in the diagram, 00 = 0x9000, 04 = 0x9004, 08 = 0x9008, etc.
- None of the data in the heap is initialized (represented by "?")
- We need to set up the heap so we can use it, so we first put in a header that tells us that the whole heap is free.

- We need to set up the heap so we can use it, so we first put in a header that tells us that the whole heap is free.

96 bytes

00	04	08	0c	10	14	18	1c	20	24	28	2c	30	34	38	3c	40	44	48	4c	50	54	58	5c
88	0	?	?	?	?	?	?	?	?	?	?	?	?	?	?	?	?	?	?	?	?	?	?

The shaded area is our header, and it says that there are **88** bytes free in the heap? Why **88**?

How can we use our struct to set this value? Assume that `heap_start` is a `void *` pointer to the start of our heap:

```
(struct header *)heap_start->payload_size = 88;
(struct header *)heap_start->status = 0;
```

In your malloc.c code, you will want to do this without the **88** constant. If you want to know the size of a header, you can simply write:

```
sizeof(struct header)
```

- Let's look at the following sets of mallocs, frees, and reallocs:

```
a = malloc(16);
b = malloc(8);
c = malloc(16);
d = malloc(8);
e = malloc(8);
free(d);
free(c);
```

	Value
a	
b	
c	
d	
e	

96 bytes

00	04	08	0c	10	14	18	1c	20	24	28	2c	30	34	38	3c	40	44	48	4c	50	54	58	5c
88	0	?	?	?	?	?	?	?	?	?	?	?	?	?	?	?	?	?	?	?	?	?	?

- Let's look at the following sets of mallocs, frees, and reallocs:

```
a = malloc(16);
b = malloc(8);
c = malloc(16);
d = malloc(8);
e = malloc(8);
free(d);
free(c);
```

	Value
a	
b	
c	
d	
e	

96 bytes

00	04	08	0c	10	14	18	1c	20	24	28	2c	30	34	38	3c	40	44	48	4c	50	54	58	5c

| 88 | 0 | ? | ? | ? | ? | ? | ? | ? | ? | ? | ? | ? | ? | ? | ? | ? | ? | ? | ? | ? | ? | ? | ? |

- Let's look at the following sets of mallocs, frees, and reallocs:

```
a = malloc(16);
b = malloc(8);
c = malloc(16);
d = malloc(8);
e = malloc(8);
free(d);
free(c);
```

	Value
a	**0x9008**
b	
c	
d	
e	

96 bytes

00	04	08	0c	10	14	18	1c	20	24	28	2c	30	34	38	3c	40	44	48	4c	50	54	58	5c
16	1	a	a	a	a	64	0	?	?	?	?	?	?	?	?	?	?	?	?	?	?	?	?

- We change the header at 0x9000 to 16 for the allocation, and set the status to 1 ("used").
- We then add another header for the next block, and set its value to 64 and the status to 0 ("free")
- We then pass back 0x9008 to the calling function.

- Let's look at the following sets of mallocs, frees, and reallocs:

```
a = malloc(16);
b = malloc(8);
c = malloc(16);
d = malloc(8);
e = malloc(8);
free(d);
free(c);
```

	Value
a	**0x9008**
b	
c	
d	
e	

96 bytes

00	04	08	0c	10	14	18	1c	20	24	28	2c	30	34	38	3c	40	44	48	4c	50	54	58	5c
16	1	a	a	a	a	64	0	?	?	?	?	?	?	?	?	?	?	?	?	?	?	?	?

- We first check the initial header at location 0x9008 and find that we cannot allocate space there (it is used), so we jump ahead by 16 bytes from the end of our header to find the next header. We find that there are 64 bytes free, enough for our allocation of 8, so we place the block there...

- Let's look at the following sets of mallocs, frees, and reallocs:

```
a = malloc(16);
b = malloc(8);
c = malloc(16);
d = malloc(8);
e = malloc(8);
free(d);
free(c);
```

	Value
a	**0x9008**
b	**0x9020**
c	
d	
e	

96 bytes

00	04	08	0c	10	14	18	1c	20	24	28	2c	30	34	38	3c	40	44	48	4c	50	54	58	5c
16	1	a	a	a	a	8	1	b	b	48	0	?	?	?	?	?	?	?	?	?	?	?	?

- We have now updated b's header, and put a header after it to designate the rest of the free block.
- We return 0x9020 to the calling function.

- Let's look at the following sets of mallocs, frees, and reallocs:

```
a = malloc(16);
b = malloc(8);
c = malloc(16);
d = malloc(8);
e = malloc(8);
free(d);
free(c);
```

	Value
a	**0x9008**
b	**0x9020**
c	
d	
e	

96 bytes

00	04	08	0c	10	14	18	1c	20	24	28	2c	30	34	38	3c	40	44	48	4c	50	54	58	5c
16	1	a	a	a	a	8	1	b	b	48	0	?	?	?	?	?	?	?	?	?	?	?	?

- We walk the heap to find a free block that fits for c, as well.

- Let's look at the following sets of mallocs, frees, and reallocs:

```
a = malloc(16);
b = malloc(8);
c = malloc(16);
d = malloc(8);
e = malloc(8);
free(d);
free(c);
```

	Value
a	**0x9008**
b	**0x9020**
c	**0x9030**
d	
e	

96 bytes

00	04	08	0c	10	14	18	1c	20	24	28	2c	30	34	38	3c	40	44	48	4c	50	54	58	5c
16	1	a	a	a	a	8	1	b	b	16	1	c	c	c	c	24	0	?	?	?	?	?	?

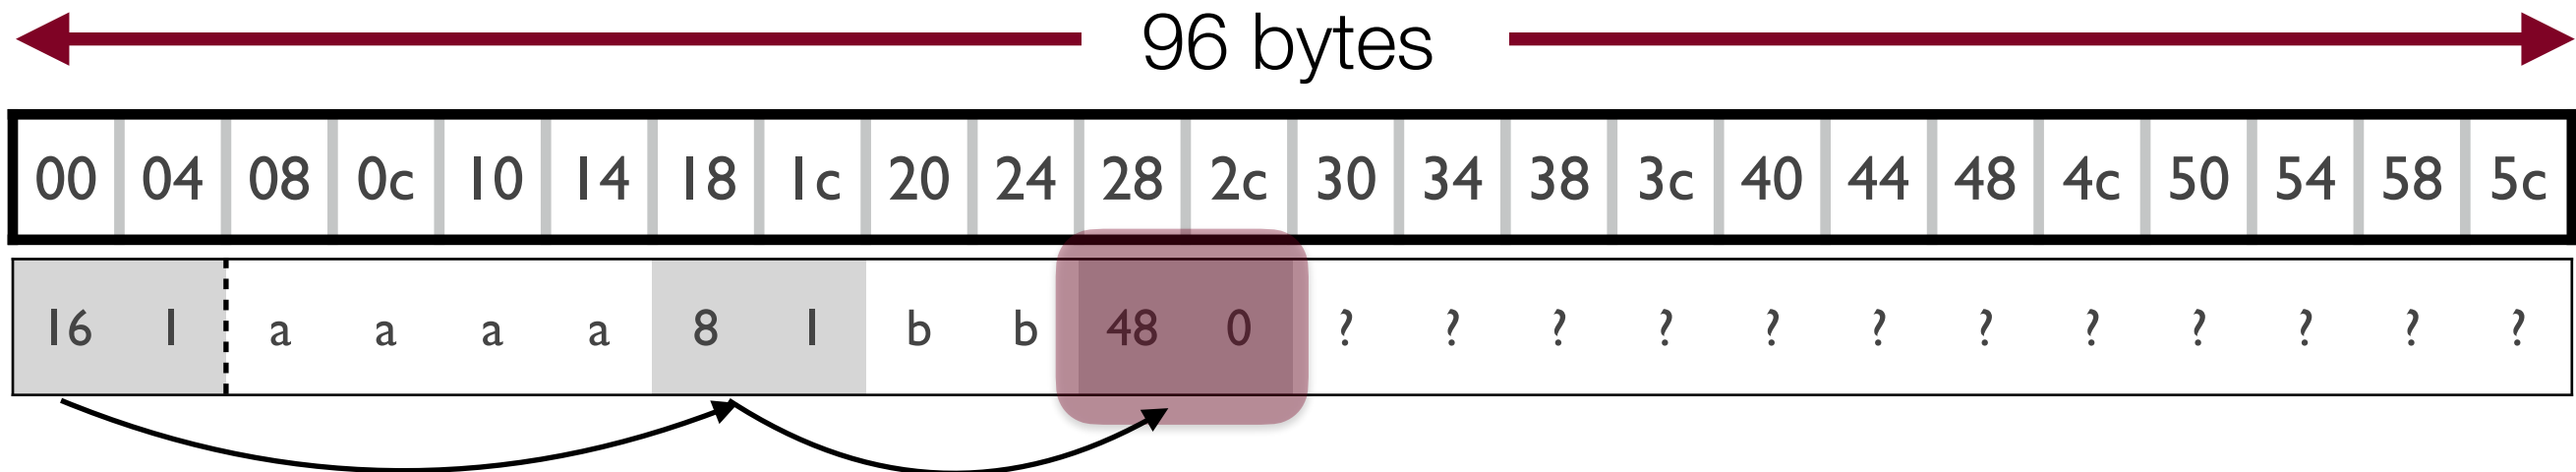- We return 0x9030 to the calling function.

- Let's look at the following sets of mallocs, frees, and reallocs:

```
a = malloc(16);
b = malloc(8);
c = malloc(16);
d = malloc(8);
e = malloc(8);
free(d);
free(c);
```

	Value
a	**0x9008**
b	**0x9020**
c	**0x9030**
d	**0x9048**
e	**0x9058**

96 bytes

00	04	08	0c	10	14	18	1c	20	24	28	2c	30	34	38	3c	40	44	48	4c	50	54	58	5c
16	1	a	a	a	a	8	1	b	b	16	1	c	c	c	c	8	1	d	d	8	1	e	e

- We repeat the process for d and e.

- Let's look at the following sets of mallocs, frees, and reallocs:

```
a = malloc(16);
b = malloc(8);
c = malloc(16);
d = malloc(8);
e = malloc(8);
free(d);
free(c);
```

|	Value
a	**0x9008**
b	**0x9020**
c	**0x9030**
d	**0x9048**
e	**0x9058**

96 bytes

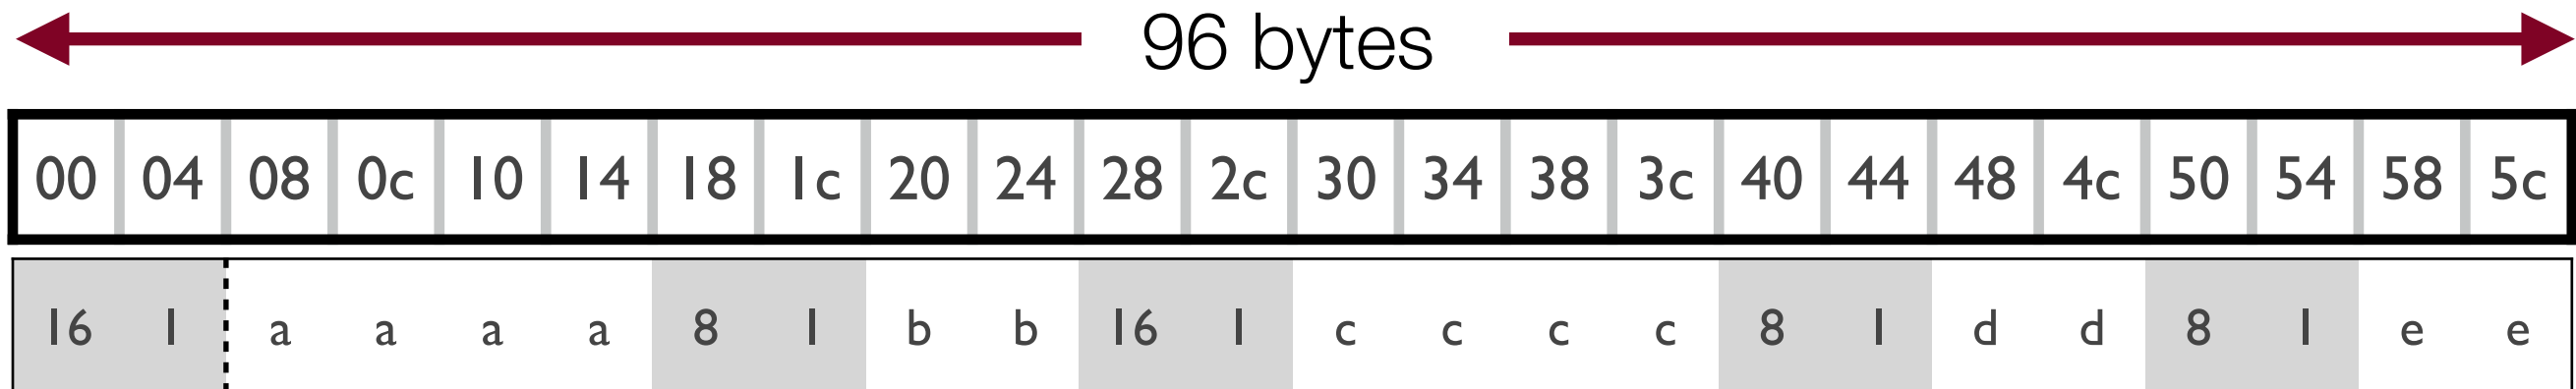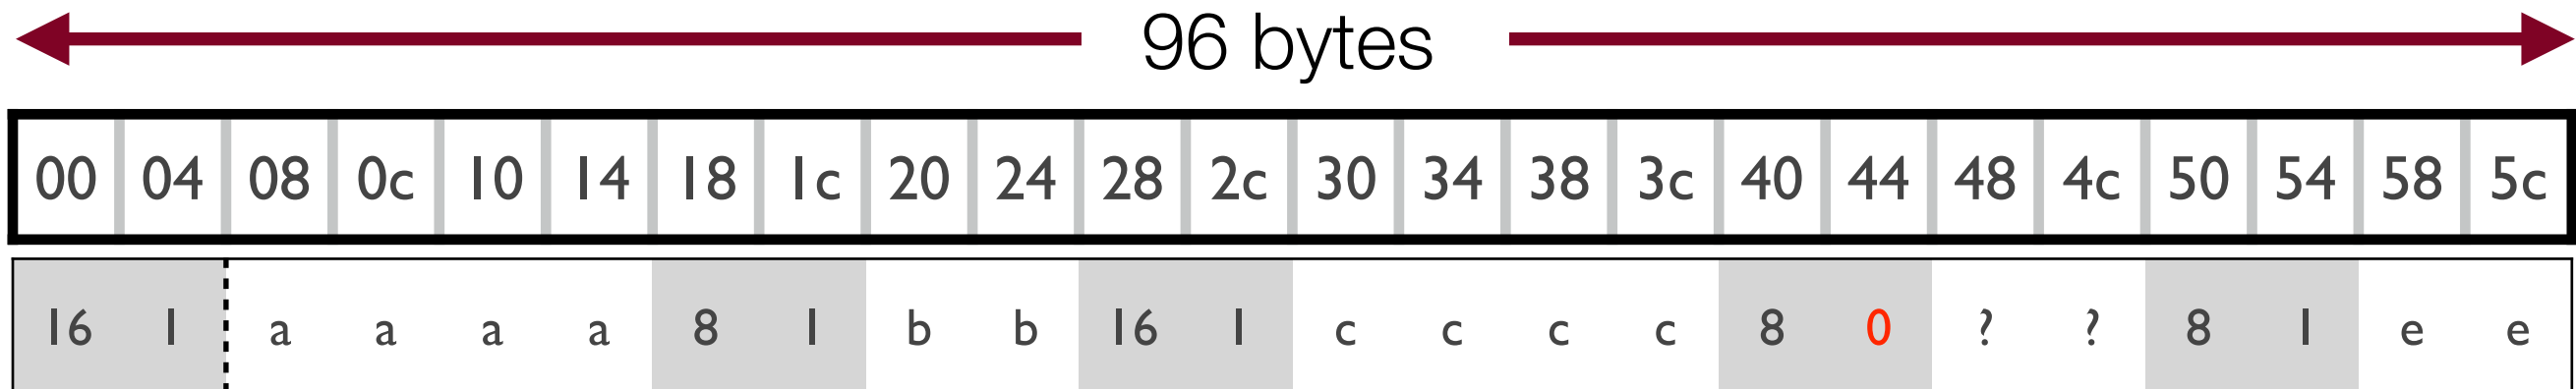| 00	04	08	0c	10	14	18	1c	20	24	28	2c	30	34	38	3c	40	44	48	4c	50	54	58	5c
16	1	a	a	a	a	8	1	b	b	16	1	c	c	c	c	8	1	d	d	8	1	e	e

- Now we need to free d.

- Let's look at the following sets of mallocs, frees, and reallocs:

```
a = malloc(16);
b = malloc(8);
c = malloc(16);
d = malloc(8);
e = malloc(8);
free(d);
free(c);
```

	Value
a	**0x9008**
b	**0x9020**
c	**0x9030**
d	
e	**0x9058**

← 96 bytes →

| 00	04	08	0c	10	14	18	1c	20	24	28	2c	30	34	38	3c	40	44	48	4c	50	54	58	5c
16	1	a	a	a	a	8	1	b	b	16	1	c	c	c	c	8	0	?	?	8	1	e	e

- Now we need to free d.
- We mark d as free, and then see if the follow-on block is free (it isn't).

- Let's look at the following sets of mallocs, frees, and reallocs:

```
a = malloc(16);
b = malloc(8);
c = malloc(16);
d = malloc(8);
e = malloc(8);
free(d);
free(c);
```

	Value
a	**0x9008**
b	**0x9020**
c	**0x9030**
d	
e	**0x9058**

96 bytes

| 00	04	08	0c	10	14	18	1c	20	24	28	2c	30	34	38	3c	40	44	48	4c	50	54	58	5c
16	1	a	a	a	a	8	1	b	b	16	1	c	c	c	c	8	0	?	?	8	1	e	e

- Now we need to free c.

- Let's look at the following sets of mallocs, frees, and reallocs:

```
a = malloc(16);
b = malloc(8);
c = malloc(16);
d = malloc(8);
e = malloc(8);
free(d);
free(c);
```

	Value
a	**0x9008**
b	**0x9020**
c	
d	
e	**0x9058**

96 bytes

00	04	08	0c	10	14	18	1c	20	24	28	2c	30	34	38	3c	40	44	48	4c	50	54	58	5c
16	1	a	a	a	a	8	1	b	b	16	0	?	?	?	?	8	0	?	?	8	1	e	e

- We mark c as free, and then check to see if there is a free block to the right -- there is!

- Let's look at the following sets of mallocs, frees, and reallocs:

```
a = malloc(16);
b = malloc(8);
c = malloc(16);
d = malloc(8);
e = malloc(8);
free(d);
free(c);
```

	Value
a	**0x9008**
b	**0x9020**
c	
d	
e	**0x9058**

← 96 bytes →

00	04	08	0c	10	14	18	1c	20	24	28	2c	30	34	38	3c	40	44	48	4c	50	54	58	5c
16	1	a	a	a	a	8	1	b	b	32	0	?	?	?	?	?	?	?	?	8	1	e	e

- We mark c as free, and then check to see if there is a free block to the right -- there is!
- We "coalesce" the blocks because they are next to each other.
- We gain back a bit of space, and we also get a bigger block to allocate later.