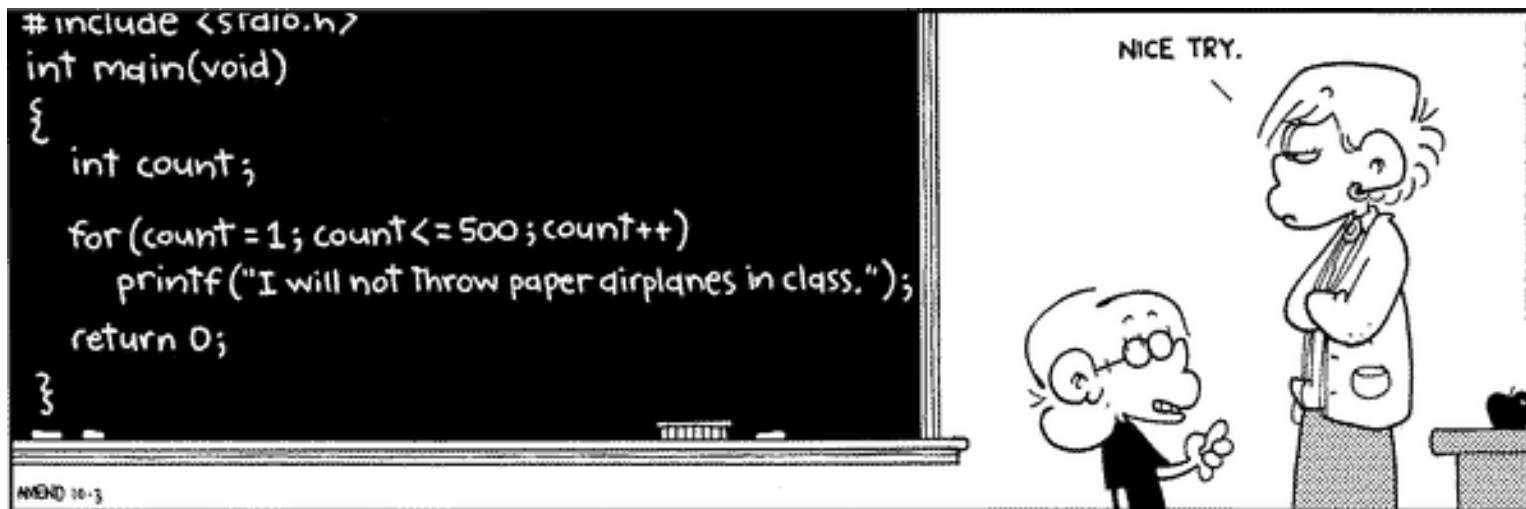


Admin

- Halfway there!
- A deep appreciation for printf -- and you'll appreciate having it!



Today: Thanks for the memory!

Linker memory map, address space layout

Loading, how an executable file becomes a running program

Heap allocation, malloc and free

**Combining Multiple Modules (.o)
into a
Single Executable (.elf)**

memmap

```
// memmap
MEMORY
{
    ram : ORIGIN = 0x8000,
        LENGTH = 0x8000000
}
.text : {
    start.o (.text)
    *(.text*)
} > ram

// Why must start.o go first?
```

**`_start` must be located
at `#0x8000!!`**

Magic constant that's part of Raspberry Pi boot sequence.

```
$ cd code/demo
$ arm-none-eabi-nm -n clock.elf
00008000 T _start
0000800c t hang
00008010 T square
0000801c T blink
00008070 T main
0000808c T timer_init
00008090 T timer_get_ticks
00008098 T timer_delay_us
000080a4 T timer_delay_ms
000080c0 T timer_delay
000080e0 T gpio_init
000080e4 T gpio_set_function
000080e8 T gpio_get_function
000080f0 T gpio_set_input
000080f4 T gpio_set_output
000080f8 T gpio_write
000080fc T gpio_read
00008104 T _cstart
00008154 T __bss_start__
00008158 T __bss_end__
```

```
# size reports the size of the text
% arm-none-eabi-size main.elf
  text    data    bss    dec    hex filename
   216      0      0    216    d8 main.elf
```

```
% arm-none-eabi-size *.o
  text    data    bss    dec    hex filename
     8      0      0      8      8 clock.o
    80      0      0    80    50 cstart.o
    20      0      0    20    14 gpio.o
    20      0      0    20    14 main.o
    12      0      0    12      c start.o
    76      0      0    76    4c timer.o
```

```
# Note that the sum of the sizes of the .o's
# is equal to the size of the main.exe
```

Relocation

```
// start.s
```

```
.globl _start
```

```
start:
```

```
    mov sp, #0x80000000
```

```
    mov fp, #0
```

```
    bl _cstart
```

```
hang:
```

```
    b hang
```



```
// Disassembly of start.o (start.list)
```

```
00000000 <_start>:
```

```
    0:    mov sp, #0x80000000
```

```
    4:    mov fp, #0
```

```
    8:    bl    0 <_cstart>
```

```
0000000c <hang>:
```

```
    c:    b     c <hang>
```

```
// Note: the address of _cstart is 0
```

```
// Why?
```

```
//    _start doesn't know where c_start is!
```

```
// Note it does know the address of hang
```

```
// Disassembly of clock.elf.list
```

```
00008000 <_start>:
```

```
    8000: mov        sp, #134217728    ; 0x80000000
```

```
    8004: bl          8088 <_cstart>
```

```
00008008 <hang>:
```

```
    8008: b          8008 <hang>
```

```
00008088 <_cstart>:
```

```
    8088:    push    {r3, lr}
```

```
// Note: the address of _cstart is #8088
```

```
// Now _start knows where _cstart is!
```

data/

```
// uninitialized global and static global  
variables
```

```
int i;  
static int j;
```

```
// initialized global and static global variables
```

```
int k = 1;  
int l = 0;  
static int m = 2;
```

```
// const global and const static global
```


```
const int n = 3;  
static const int o = 4;
```

Function Static?

A static variable declared in a function exists for the lifetime of the program: it's a like a global variable in this way

However, it is *only accessible within that scope*, so it's not a global variable

```
int function(int x) {  
    static int counter = 0;  
    counter++;  
    return x + counter;  
}
```



counter lives like a global variable, but can only be named within function

```
% arm-none-eabi-nm -S tricky.o
000000004 000000004 C i
000000000 000000004 b j
000000000 000000004 D k
000000004 000000004 B l
000000004 000000004 d m
000000000 000000004 R n
000000004 000000004 r o
000000000 0000000d8 T tricky
```

```
# The global uninitialized variable i
# is in common (C).
```

```
# If you compile with -Og, some variables
# are optimized out -- which?
```

Guide to Symbols

T/t - text

D/d - read-write data

R/r - read-only data

B/b - bss (*Block Started by Symbol*)

C - common (instead of B)

lower-case letter means `static`

Data Symbols

Types

- global vs static
- read-only data vs data
- initialized vs uninitialized data
- common (shared data)

Sections

Instructions go in `.text`

Data goes in `.data`

const data (read-only) goes in `.rodata`

Uninitialized data goes in `.bss`

+ other information about the program

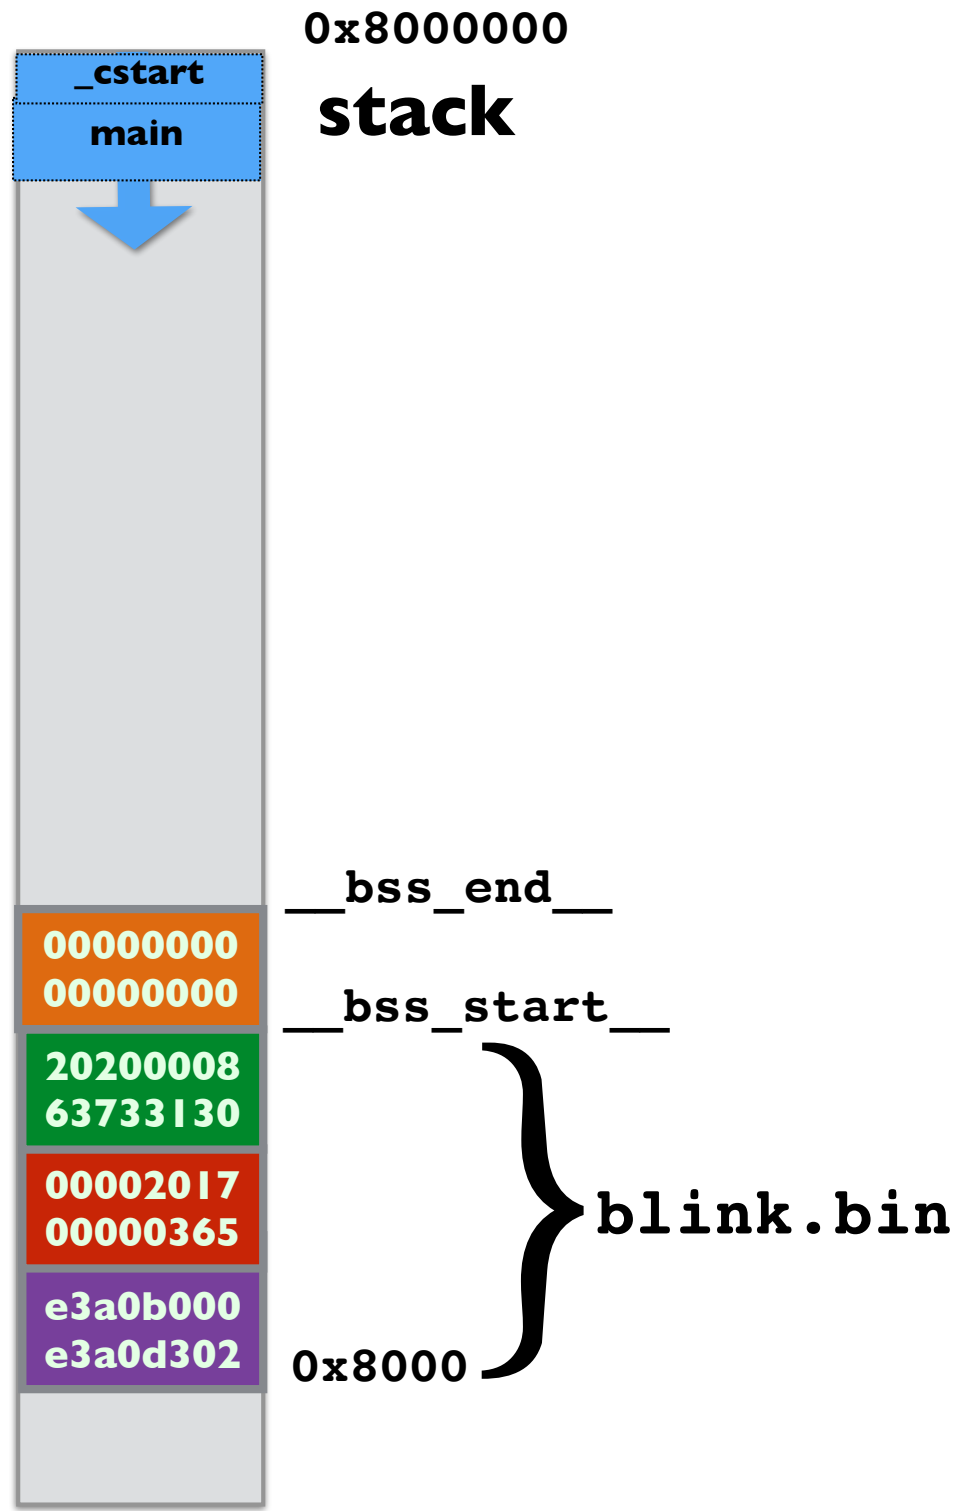
- symbols, relocation, debugging, ...

SECTIONS

```
{
    .text 0x8000 : { start.o(.text*)
                    *(.text*) }
    .data :      { *(.data*) }
    .rodata :    { *(.rodata*) }

    __bss_start__ = .;
    .bss :        { *(.bss*)
                    *(COMMON) }
    __bss_end__ = ALIGN(8);
}
```

(zeroed data) **.bss**
(read-only data) **.rodata**
(initialized data) **.data**
.text



```
$ cd code/data
$ arm-none-eabi-nm -S -n main.elf
00008000 T _start
0000800c t hang
00008010 00000038 T main
00008048 00000040 T tricky
00008088 00000058 T _cstart
000080e0 00000004 D k
000080e4 00000004 R n
000080e8 R __bss_start__
000080e8 00000004 b j
000080ec 00000004 B l
000080f0 00000004 B i
000080f8 B __bss_end__
```

SECTIONS

```
{  
    .text 0x8000 : {  
        start.o(.text*)  
        *(.text*)  
    }  
    .data : { *(.data*) }  
    .rodata : { *(.rodata*) }  
    __bss_start__ = .;  
    .bss : { *(.bss*) *(COMMON) }  
    __bss_end__ = ALIGN(8);  
}
```

SECTIONS

```
{
    .text 0x8000 : { start.o(.text*)
                    *(.text*)}

    .data :        { *(.data*) }
    .rodata :      { *(.rodata*) }

    __bss_start__ = .;
    .bss :         { *(.bss*)
                    *(COMMON) }
    __bss_end__ = ALIGN(8);
}
```

Use this memory for heap 

(zeroed data) **.bss**

(read-only data) **.rodata**

(initialized data) **.data**

.text



0x8000000

```
_start:
    mov sp, #0x8000000
    mov fp, #0
    bl _cstart
```

```
void _cstart(void) {
    int *bss = &__bss_start__;
    while (bss < &__bss_end__)
        *bss++ = 0;
}

main();
}
```

`__bss_end__`

`__bss_start__`

blink.bin

0x8000

Global allocation

- + **Convenient**

 - Fixed location, shared across entire program

- + **Fast, plentiful**

 - No explicit allocation/deallocation

 - But have to send over serial to bootloader (can be slow)

- **Size fixed at declaration, no option to resize**

- +/- **Scope and lifetime is global**

 - No encapsulation, hard to track use/dependencies

 - One shared namespace, have to manually manage conflicts

 - Static variables can address some issues

 - Frowned upon stylistically (advanced systems reasons)

Stack allocation

- + **Convenient**

 - Automatic alloc/dealloc on function entry/exit

- + **Fast**

 - Fast to allocate/deallocate, good locality

- **Usually don't allocate large chunks (megabytes)**

- **Size fixed at declaration, no option to resize**

- +/- **Scope/lifetime dictated by control flow**

 - Private to stack frame

 - Does not persist after function exits

- **Memory bug can corrupt execution**

Heap allocation

- + **Moderately efficient**

 - Have to search for available space, update record-keeping

- + **Very plentiful**

 - Heap enlarges on demand to limits of address space

- + **Versatile, under programmer control**

 - Can precisely determine scope, lifetime

 - Can be resized

- **Low type safety (can't access by value)**

 - Interface is raw void *, number of bytes

- **Lots of opportunity for error**

 - (allocate wrong size, use after free, double free)

- **Leaks**

- **Hard to track down sources of corruption**

Heap interface

```
void *malloc (size_t nbytes);  
void free (void *ptr);  
void *realloc (void *ptr, size_t nbytes);
```

void* pointer

"Generic" pointer, a memory address

Type of pointee is not specified, unknown

What you can do with a void*

Pass to/from function, pointer assignment

What you cannot do with a void*

Cannot dereference (must cast first)

Cannot do pointer arithmetic (cast to char * to manually control scaling)

Why do we need a heap?

Let's see an example!

code/heap/names.c

How to implement a heap

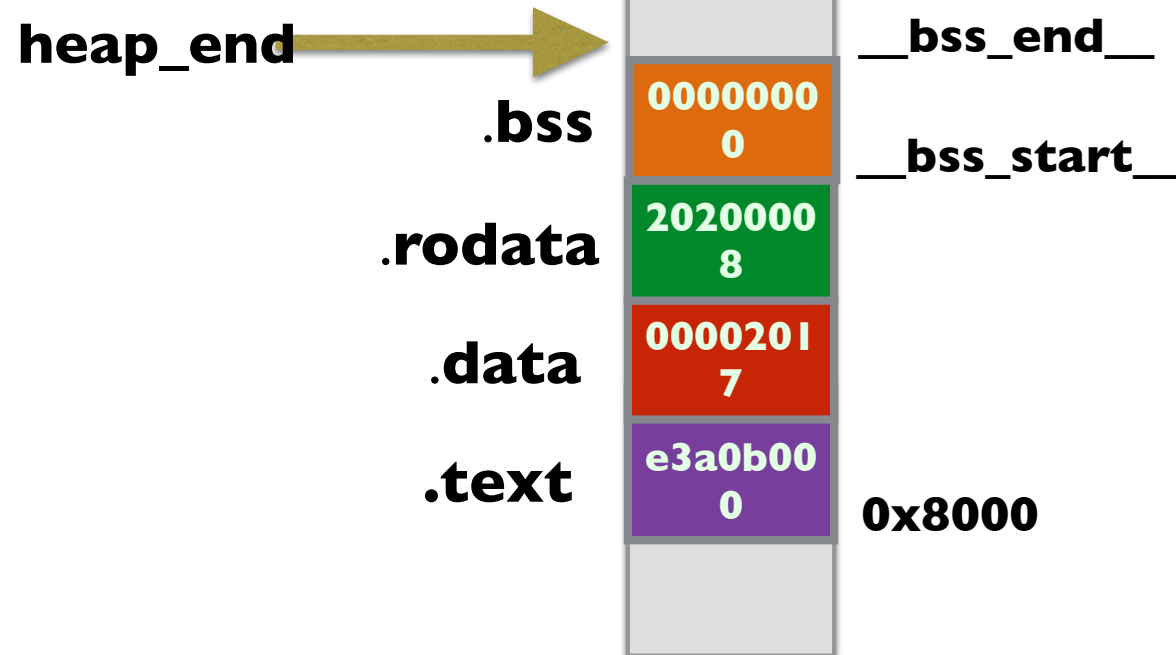


```

void *sbrk(int nbytes)
{
    static void *heap_end = &__bss_end__;

    void *prev_end = heap_end;
    heap_end = (char *)heap_end + nbytes;
    return prev_end;
}


```




Tracing the bump allocator

Each square represents 8 bytes of space


p1 = malloc(32)




p2 = malloc(40)




p3 = malloc(48)



free(p2)



p4 = malloc(16)



Bump Memory Allocator

`code/heap/malloc.c`


Evaluate bump allocator

- + Operations super-fast
- + Very simple code, easy to verify, test, debug
- No recycling/re-use
 - (in what situations will this be problematic?)
- Sad consequences when `sbrk()` advances into stack
 - (what can we do about that?)


Pre-block header, implicit list

Each square represents 8 bytes of space, size recorded as total byte count


p1 = malloc(32)




p2 = malloc(40)




p3 = malloc(48)



free(p2)



p4 = malloc(16)



Header struct

```
struct header {
    unsigned int size;
    unsigned int status;
};                                     // sizeof(struct header) = 8 bytes

enum { IN_USE = 0, FREE = 1};

void *malloc(size_t nbytes)
{
    nbytes = roundup(nbytes, 8);
    size_t total_bytes = nbytes + sizeof(struct header);

    struct header *hdr = sbrk(total_bytes);
    hdr->size = nbytes;
    hdr->status = IN_USE;
    return hdr + 1;    // return address at start of payload
}
```

Challenges for malloc client

- **Correct allocation (size in bytes)**
- **Correct access to block (within bounds, not freed)**
- **Correct free (once and only once, at correct time)**

What happens if you...

- forget to free a block after you are done using it?
- access a memory block after you freed it?
- free a block twice?
- free a pointer you didn't malloc?
- access outside the bounds of a heap-allocated block?

Challenges for malloc implementor

just **malloc** is easy 😎

malloc with **free** is hard 🤔

Efficient **malloc** with **free**Yikes! 😓

Complex code (pointer math, typecasts)

Thorough testing is challenge (more so than usual)

Critical system component

correctness is non-negotiable, ideally fast and compact

Survival strategies:

draw pictures

printf (you've earned it!!)

early tests use examples small enough to trace by hand if need be

build up to bigger, more complex tests