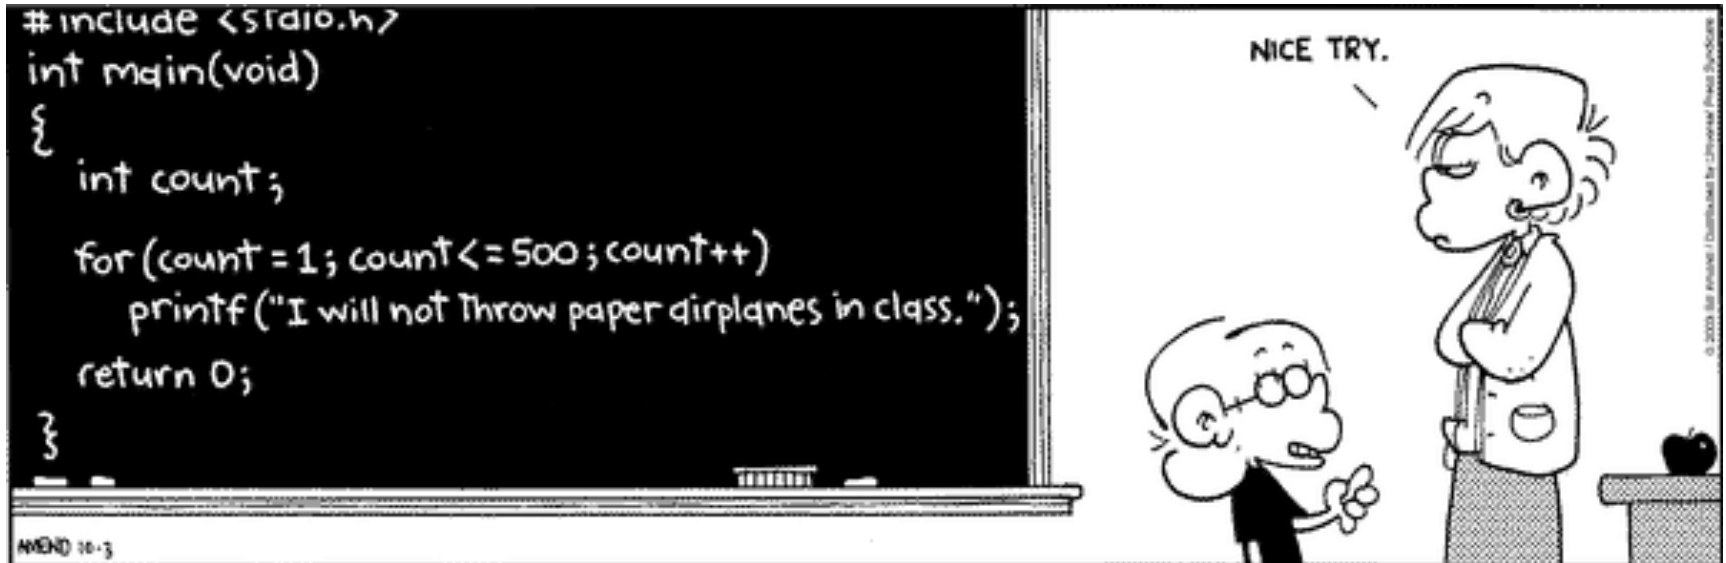


Admin

- Assign 2 issues posted, open for fix and resubmit
- printf perseverance and pride!!



Today: Thanks for the memory!

Linker memory map, address space layout

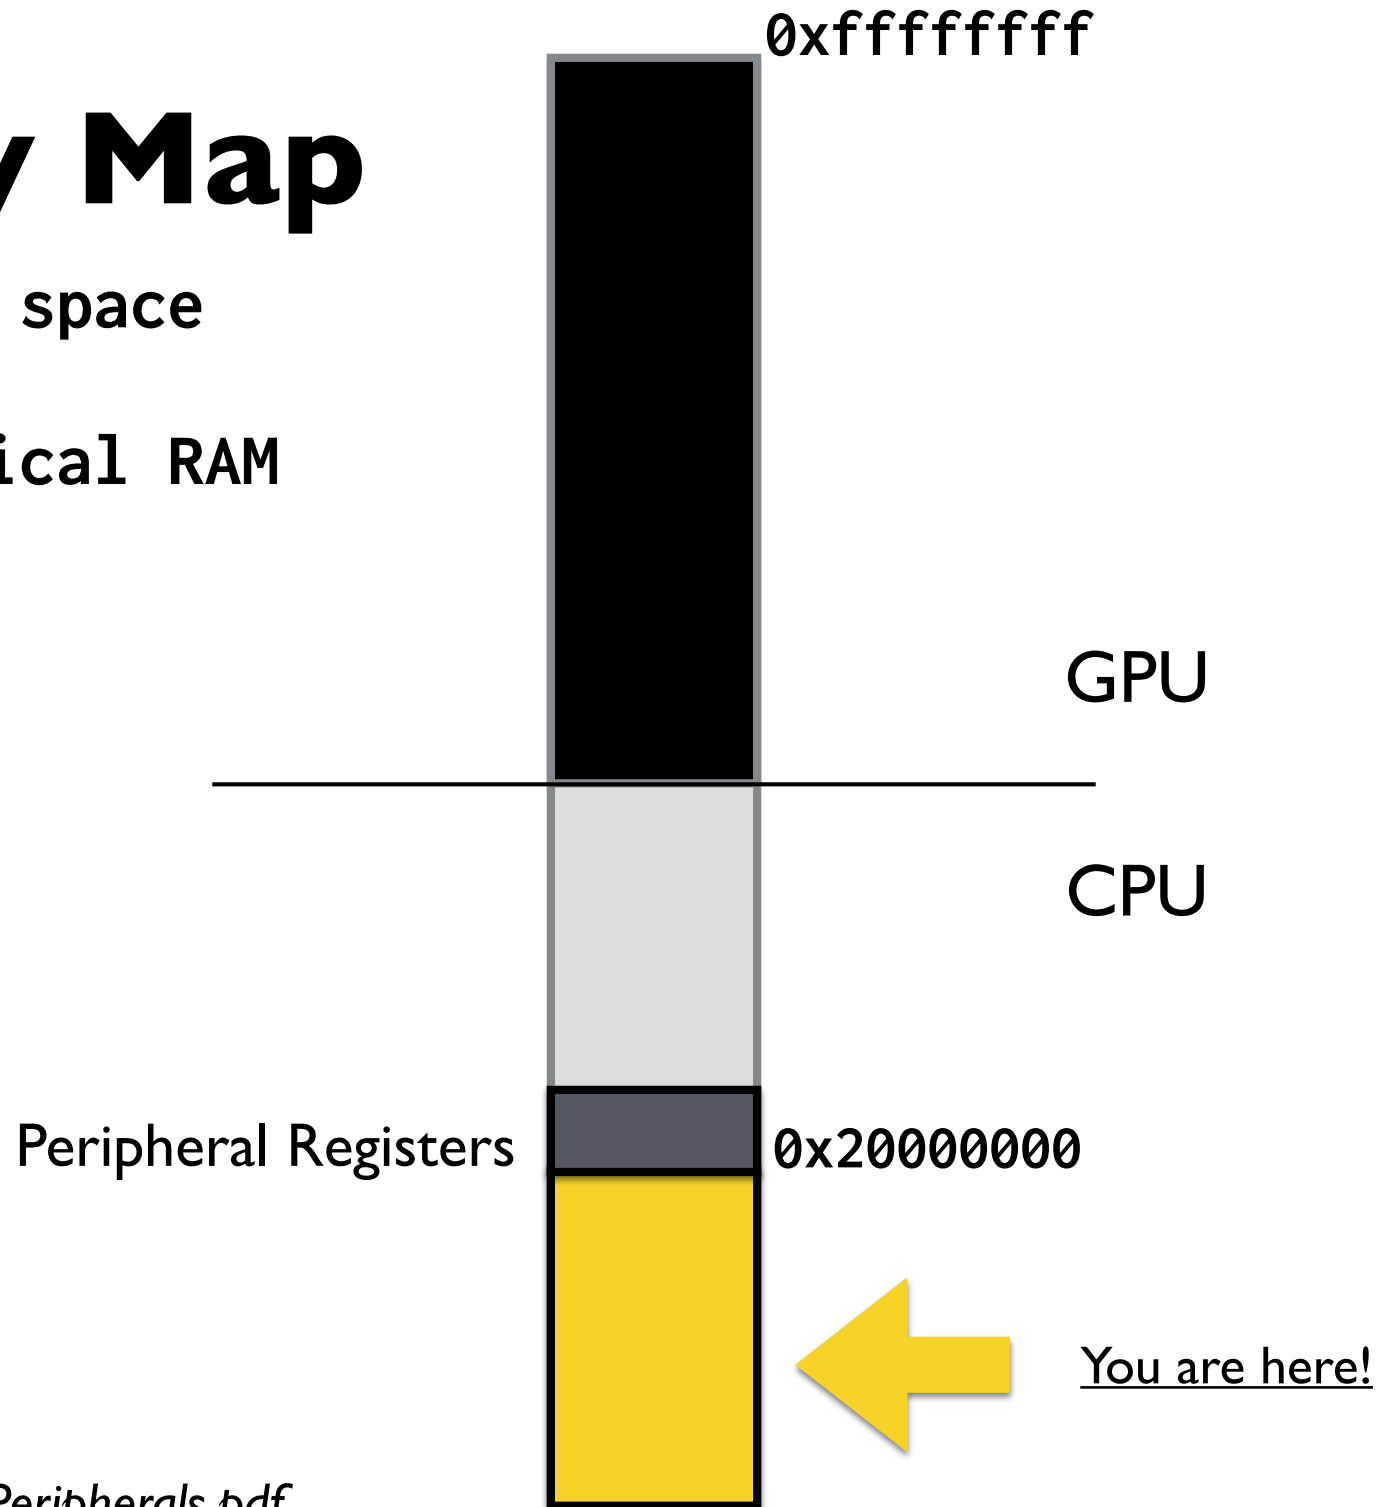
Loading, how an executable file becomes a running program

Heap allocation, malloc and free

Memory Map

32-bit address space

512 MB of physical RAM



SECTIONS

```
{
    .text 0x8000 : { *(<div data-bbox="545 12 645 960" data-label="Diagram">



The diagram illustrates the memory layout of a program. It shows a vertical stack of sections. At the top, there is a grey section labeled _cstart and a blue section labeled main. Below these is a large grey section. Further down is a white section containing a cartoon character, representing the .bss section. Below that are four colored sections: orange for .bss, green for .data, red for .rodata, and purple for .text. At the bottom is another grey section. Arrows indicate the stack growing downwards from _cstart/main and the heap growing upwards from the .bss section. A large bracket on the right groups the .bss, .data, .rodata, and .text sections, labeling them as blink.bin. Address values are provided for each section: 0x8000000 for _cstart, 00000000 for .bss, 20200008 for .data, 00002017 for .rodata, and e3a0b000 for .text. A size of 0x8000 is indicated for the .text section.


```

Global allocation

- + **Convenient**

 - Fixed location, shared across entire program

- + **Fairly efficient, plentiful**

 - No explicit allocation/deallocation

 - Heavy cost to send over serial line to bootloader

- + **Reasonable type safety**

- **Size fixed at declaration, no option to resize**

- +/- **Scope and lifetime is global**

 - No encapsulation, hard to track use/dependencies

 - One shared namespace, have to manually manage conflicts

 - Frowned upon stylistically

Stack allocation

- + **Convenient**

 - Automatic alloc/dealloc on function entry/exit

- + **Efficient, fairly plentiful**

 - Fast to allocate/deallocate, low consequence if large size

- + **Reasonable type safety**

- **Size fixed at declaration, no option to resize**

- +/- **Scope/lifetime dictated by control flow**

 - Private to stack frame

 - Does not persist after function exits

Heap allocation

- + **Moderately efficient**

 - Have to search for available space, update record-keeping

- + **Very plentiful**

 - Heap enlarges on demand to limits of address space

- + **Versatile, under programmer control**

 - Can precisely determine scope, lifetime

 - Can be resized

- **Low type safety**

 - Interface is raw void *, number of bytes

- **Much opportunity for error**

 - (allocate wrong size, use after free, double free)

- **Leaks** (less critical, but annoying nonetheless)

Heap interface

```
void *malloc (size_t nbytes);  
void free (void *ptr);
```

void* pointer

"Generic" pointer, a memory address

Type of pointee is not specified, could be any data

What you can do with a void*

Pass to/from function, pointer assignment

What you cannot do with a void*

Cannot dereference (must cast first)

Cannot do pointer arithmetic (cast to char * to manually control scaling)

Cannot use array indexing (size of pointee not known!)

Why do we need a heap?

Let's see an example!

`code/heap/names.c`

How to implement a heap

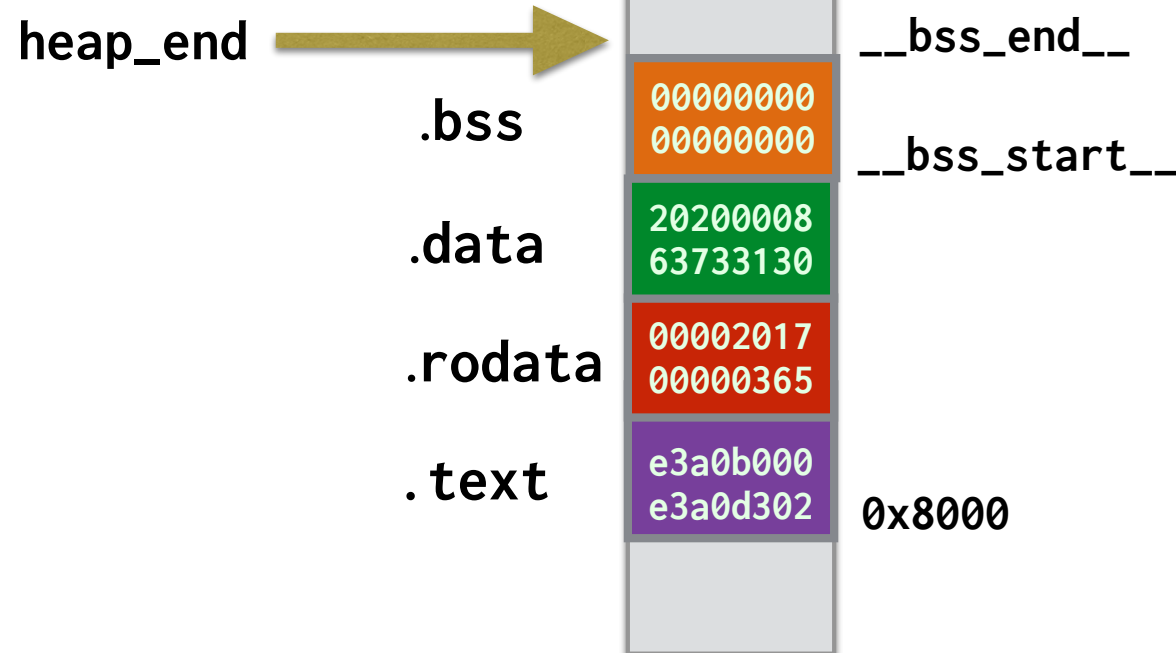


```

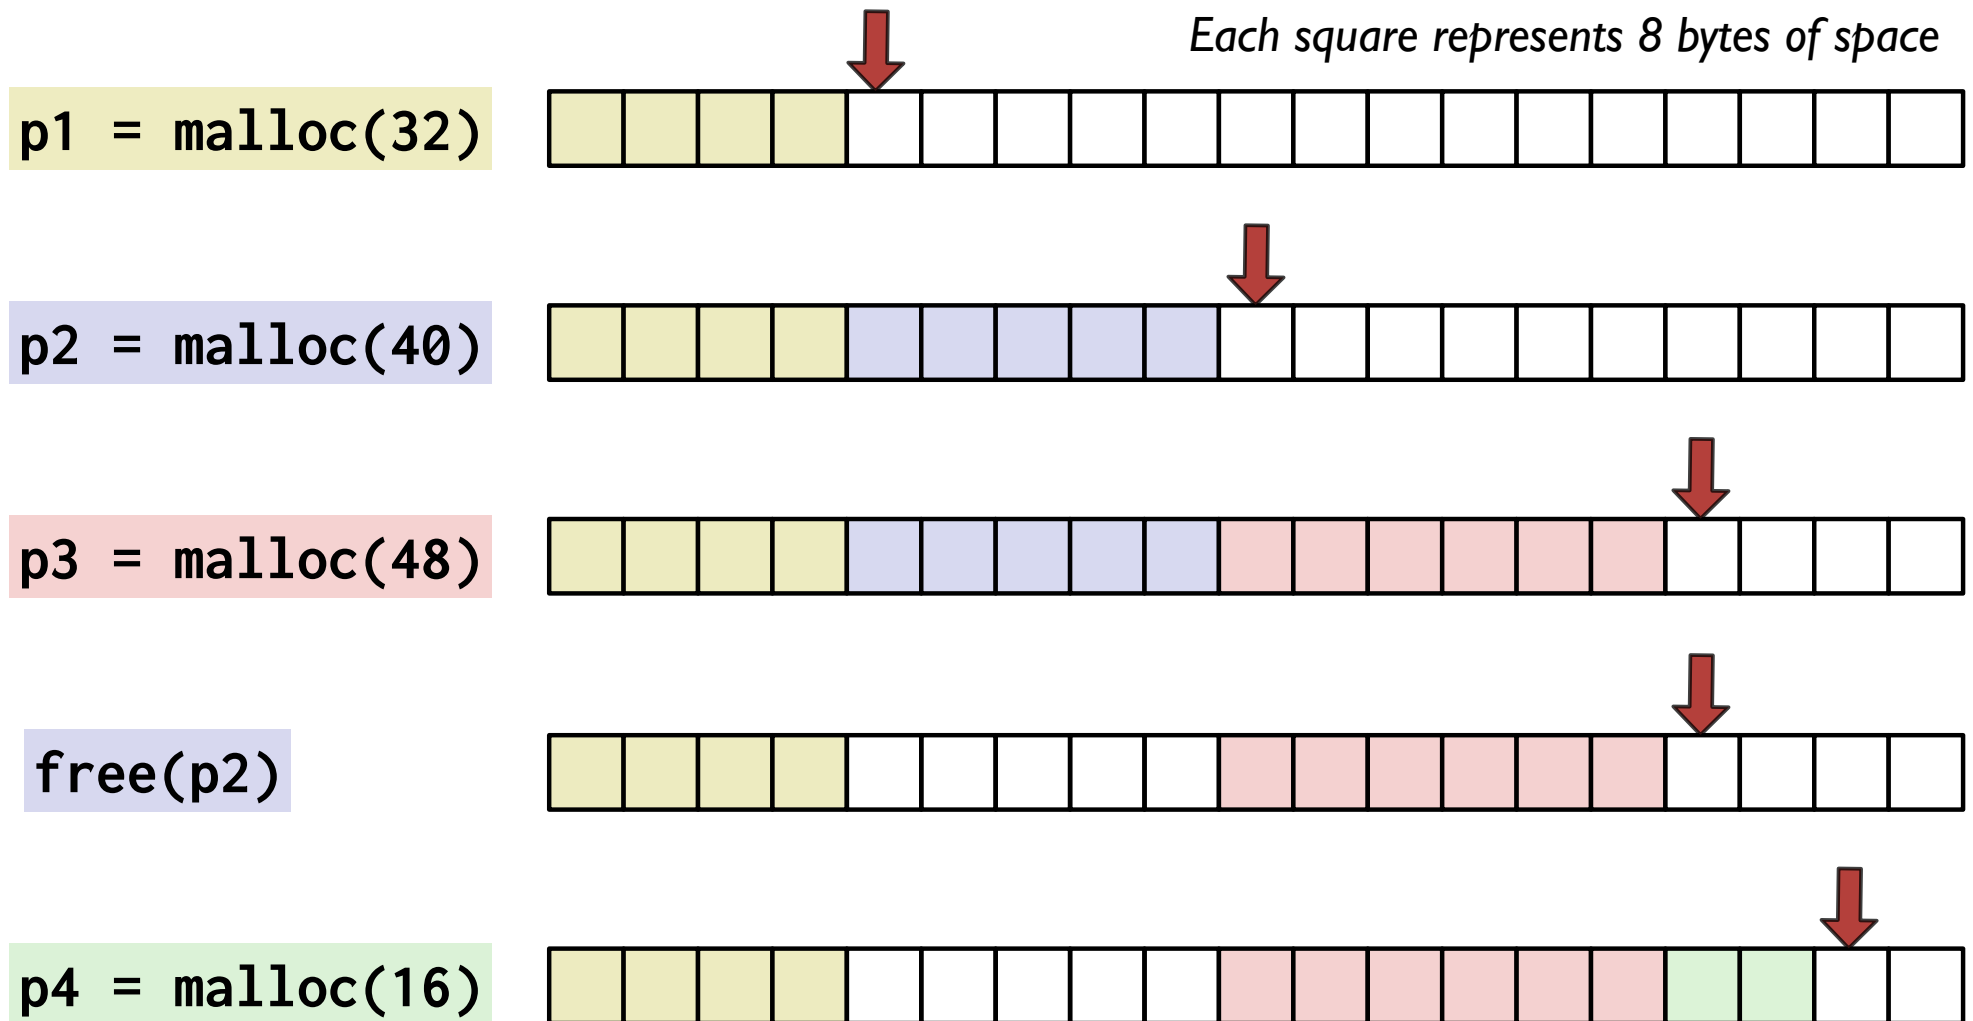
void *sbrk(int nbytes)
{
    static void *heap_end = &__bss_end__;

    void *prev_end = heap_end;
    heap_end = (char *)heap_end + nbytes;
    return prev_end;
}

```



Tracing the bump allocator



Bump Memory Allocator

`code/heap/malloc.c`


Evaluate bump allocator

- + Operations super-fast
- + Very simple code, easy to verify, test, debug
- No recycling/re-use
 - (in what situations will this be problematic?)
- Sad consequences when `sbrk()` advances into stack
 - (what can we do about that?)


Pre-block header, implicit list

Each square represents 8 bytes, header records size of payload in bytes


`p1 = malloc(32)`




`p2 = malloc(40)`




`p3 = malloc(48)`



`free(p2)`



`p4 = malloc(16)`



Header struct

```
struct header {
    unsigned int size;
    unsigned int status;
};                                     // sizeof(struct header) = 8 bytes

enum { IN_USE = 0, FREE = 1};

void *malloc(size_t nbytes)
{
    nbytes = roundup(nbytes, 8);
    size_t total_bytes = nbytes + sizeof(struct header);

    struct header *hdr = sbrk(total_bytes); // extend end of heap
    hdr->size = nbytes;
    hdr->status = IN_USE;
    return hdr + 1;    // return address at start of payload
}
```

Challenges for malloc client

- **Correct allocation (size in bytes)**
- **Correct access to block (within bounds, not freed)**
- **Correct free (once and only once, at correct time)**

What happens if you...

- forget to free a block after you are done using it?
- access a memory block after you freed it?
- free a block twice?
- free a pointer you didn't malloc?
- access outside the bounds of a heap-allocated block?

Challenges for malloc implementor

just malloc is easy 😎

malloc with free is hard 🤔

Efficient malloc with freeYikes! 😰

Complex code (pointer math, typecasts)

Thorough testing is challenge (more so than usual)

Critical system component

correctness is non-negotiable, ideally fast and compact

Survival strategies:

draw pictures

printf (you've earned it!!)

early tests use examples small enough to trace by hand if need be

build up to bigger, more complex tests